

The Proximity Prize Negative-Endpoint Ledger: Per-Field Assessments, Rescue Paths, and Sub- Lemma Close-Outs

Ömer Demirel — Snowfall Finance — omer@progfi.xyz

Proximity Prize — research ledger / companion note. **Date:** 2026-06-10. **Convention:** ABF δ -radius throughout (δ = relative proximity *radius*; $\rho = k/n$; Johnson radius $J = 1 - \sqrt{\rho}$; Singleton/MDS capacity $1 - \rho$; list-decoding capacity radius $R_{\text{cap}}(q, \rho) := H_q^{-1}(1 - \rho)$ — see §1).

0. What this ledger is

This companion note records **research-state assessments at mixed maturity levels; the status tags are load-bearing**. It is *not* a polished paper: rows below range from theorem-grade citations of the literature, through ASSEMBLED lemmas (cited components plus exact threshold arithmetic), through ESTABLISHED-MODULO-FORMALIZATION computational verdicts, down to INFERRED arithmetic flagged VERIFY and frankly OPEN items. Every claim carries its tag; a claim must never be quoted above its tag.

The polished companion is the short paper **technical-note-short.md** (same repository): the two grand challenges and their non-equivalence, the budget trichotomy in $B = 2^{-128}|F|$, the proven list-size bracket $[J - o(1), R_{\text{cap}}]$ for $|F| > 2^{128}$, the BCHKS Thm 1.9 separation, the field-agnosticity of the negative mechanism, and the open problems (keystone P'). Definitions (CA/MCA, interleaved codes, lists, the relation $(\star)$: $|\Lambda(C, \delta)| \leq |\Lambda(C^{\equiv m}, \delta)| \leq |\Lambda(C, \delta)|^m$) are as in the short paper §2. This ledger holds everything the short paper points to: the assembled per-field negative-ceiling **assessment** (internal revision rounds R13–R18), the 31/64/128/256-bit arithmetic and tables, the M31/BabyBear/circle-domain story, the 256-bit rescue-path assessments, the N1/N2 close-outs with their caveats, the experimental data, and the P' route summary.

Status-tag glossary. **PROVEN** — a theorem in the cited literature, used inside its stated range. **PROVEN BRACKET** — composition of PROVEN endpoints. **ASSEMBLED** — assembled from cited components plus exact threshold arithmetic; named soft spots listed. **ESTABLISHED-MODULO-FORMALIZATION (EMF)** — exact computation plus a structural argument, not yet written as a theorem; fail-closed. **CONDITIONALLY ESTABLISHED** — proven modulo one named, explicitly stated extension of a cited lemma. **OUR-DERIVED** — our composition/arithmetic on cited inputs. **INFERRED (VERIFY)** — arithmetic inferred from a construction's parameter window; explicitly *not* a proven endpoint. **EXPERIMENTAL** — exact small-field enumeration, corroborative only. **CONJECTURED / OPEN** — what they say.

1. Conventions

The ceiling convention (R_{cap} , load-bearing). Two nearby radii appear in the literature and in our tooling: the **exact inverse-entropy crossing** $R_{\text{cap}}(q, \rho) := H_q^{-1}(1 - \rho)$ — the radius δ solving $H_q(\delta) = 1 - \rho$, the convention of ledger §3b / `calculator/proximity_parameters.py` — and the **large- q approximation** $1 - H_q(\rho)$. **Every proven upper-ceiling statement in this ledger is a statement about R_{cap}** ; the large- q approximation appears only where explicitly labeled as such. The two are **not equal at any deployed rate** (verified by direct computation): the gap is ≤ 0.0017 across deployed rates and field sizes (maximum ≈ 0.00166 at 31-bit, $\rho = \frac{1}{8}$), with rate-dependent sign — the approximation sits *below* R_{cap} at $\rho = \frac{1}{2}$ (by $+9.6 \cdot 10^{-5}$ at 31-bit, $+1.7 \cdot 10^{-7}$ at 256-bit) and *above* it at $\rho \leq \frac{1}{4}$ (by $\approx 2.0\text{--}2.3 \cdot 10^{-5}$ at 256-bit) — i.e. the two differ **beyond the fifth decimal** at some rates even at 256-bit, and the approximation can overstate the proven ceiling, which is exactly why R_{cap} is the proven object. The large- q approximation sits a distance $H_q(\rho) - \rho \leq 1/\log_2 q$ below Singleton capacity (CS25 Claim 1; $\leq 1/31$ for 31-bit fields).

A historical software caveat — now FIXED (R29/M2 follow-up). The resolver `calculator/listsize_resolution.py` formerly capped the alphabet at `_QCAP_LOG2 = 200` before evaluating H_q , so its generated “256-bit” ceiling entries were in fact **200-bit** values — e.g. 0.49500 at $\rho = \frac{1}{2}$ vs the true 256-bit large- q -approximation value 0.49609. The cap is now 320 (no deployed field is capped) and the resolver reports the proven ceiling at the field’s true q ; all 256-bit ceiling values quoted in this ledger are true $q = 2^{256}$ values, matching `problem-ledger.md` §3b and the regenerated `listsize-analysis.md` §8. The exact 256-bit crossings R_{cap} are 0.49609/0.74681/0.87285/0.93616 at $\rho = \frac{1}{2}/\frac{1}{4}/\frac{1}{8}/\frac{1}{16}$ (the large- q approximation would give 0.49609/0.74683/0.87288/0.93618).

The notation clash (carry the ABF convention). ABF/CS25/GG25 use δ for the proximity *radius*; BCHKS and CGHLL use δ for the **code’s minimum distance** ($= 1 - \rho$, MDS) and γ (resp. θ) for the radius. BCHKS’s soundness error is a/q where $a = \#\{z : \Delta(f + zg, C) \leq \gamma\}$ and $q = |F|$; its “proximity loss” is ε^* (the gap between hypothesis and conclusion radii). When importing BCHKS into ABF: radius $\delta \equiv \gamma$, rate $\rho \equiv 1 - \delta_{\text{BCHKS}}$, error a/q with a a count. **CS25’s “capacity” is the list-decoding capacity** (our R_{cap} , which they state via the large- q approximation), strictly below the Singleton $1 - \rho$ that ABF/BCHKS/GG call capacity. **CGHLL’s “Elias radius” is likewise the large- q approximation $1 - H_q(\rho)$** ; wherever we quote their conjectures we write R_{cap} , the difference (≤ 0.0017) being below the resolution of every claim involved. All statements below are in the ABF convention.

2. The MCA bracket in detail (sub-problem 1)

This section records the *bracket* $[\delta_{\text{known-positive}}, \delta_{\text{known-negative}}]$ for the MCA challenge; the true δ_C^* lies inside it, and pinning it is open. **This is a bracket, not a theorem-complete determination.** Scoping rule (carried from the short paper §5, load-bearing): the known bracket *endpoints* are field-type insensitive; the negative *mechanism* is field-agnostic (§3.2); the *positive* beyond-Johnson program is **not** field-agnostic (§3.5).

2.1 Positive side — up to the Johnson radius (all fields)

[PROVEN] (Bordage–Chiesa–Guan–Manzur, eprint 2025/2051, Thm. 9.2). For $C = \text{RS}[F, D, k]$ over *any* domain D (including smooth subgroups), the affine-line / univariate-powers generator, message degree $d = k - 1$, and integer $m \geq 3$:

$$\varepsilon_{\text{mca}}(C, \delta) \leq \frac{(m + \frac{1}{2})^7 n^2 d}{3 \rho^{3/2} |F|}, \quad \text{valid for } \delta \leq 1 - (1 + \frac{1}{2m})\sqrt{\rho}.$$

Loss-free (= ABF ε_{mca}); as $m \rightarrow \infty$ the validity window $\rightarrow J = 1 - \sqrt{\rho}$, but larger m inflates the error via $(m + \frac{1}{2})^7$ (the m -vs- δ trade).

[PROVEN] (BCHKS, eprint 2025/2055, Thm. 1.5 = ABF Thm. 4.12 core). With $\eta = 1 - \sqrt{\rho} - \delta > 0$:

$$\varepsilon_{\text{mca}}(C, \delta) \leq \frac{C_\rho n}{\eta^5 |F|}, \quad \text{valid for } \delta \leq 1 - \sqrt{\rho} - \eta, \quad \varepsilon^* = 0.$$

The η -exponent 5 is firm (BCHKS Table 1); the leading constant C_ρ is an unspecified $O_\rho(\cdot)$ — [OUR-DERIVED caveat] we treat it as a VERIFY placeholder ($C_\rho = 1$ in the calculator). Linear in n (vs $n^2 d \sim n^3$ for Bordage–Chiesa), hence asymptotically stronger, but kept unverified until C_ρ is pinned.

Field-size consequence [OUR-DERIVED]. The *verified* Bordage–Chiesa numerator $n^2 d$ costs $\approx 2 \log_2 n + \log_2(\rho n)$ bits (≈ 60 at $n = 2^{20}$, ≈ 90 at 2^{30}). Certifying $\varepsilon_{\text{mca}} \leq 2^{-128}$ from this single-code bound therefore needs $\log_2 |F| \gtrsim 256$ for n up to 2^{30} : M31/BabyBear/Goldilocks **cannot** certify 2^{-128} from the single-code MCA bound at all (see `calculator/out/delta_star_tables.md`). Those deployments reach 128-bit soundness via the query term $(1 - \delta)^t$ plus repetition/PoW and by running the soundness argument in a large extension — not from a single-code 2^{-128} MCA certificate. (BCHKS Thm 1.5’s linear- n bound would relax this *if* C_ρ were pinned.)

$\Rightarrow \delta_{\text{known-positive}} = J = 1 - \sqrt{\rho}$ for smooth RS over all fields — **as a radius-level theorem** ($\varepsilon_{\text{mca}} \leq \text{poly}(n) d / |F|$ up to J). Whether that error meets the prize target $\varepsilon^* = 2^{-128}$ is a *separate, field-size* question: per the consequence above, only ≥ 256 -bit fields certify it from the single-code bound. Throughout this ledger, “positive to Johnson” at 31/64/128-bit fields means the radius-level theorem, **not** a certified $\varepsilon^* = 2^{-128}$ lower bound.

2.2 Negative side — near capacity (the mechanism is field-agnostic: prime and extension)

[PROVEN] (Crites–Stewart, eprint 2025/2046, Cor. 1 = ABF Thm. 4.17). For $q \geq 10$, $q \geq n$, in a k -window reaching down to the list-decoding capacity, there exist $u^{(0)}, u^{(1)}$ with $\Delta(u^{(1)}, C) > \delta$ yet $\Delta(u^{(0)} + \lambda u^{(1)}, C) \leq \delta$ for every $\lambda \in F$. Thus $\varepsilon_{\text{ca}} = 1$ reaching down to the list-decoding capacity radius R_{cap} ; δ_C^* is provably strictly below Singleton capacity.

[PROVEN] (Kambiré, arXiv:2604.09724, Thm. 1; = the KK25 / CGHLL Thm. 37 construction). For every $C > 0$ and rate $\rho \in (0, \frac{1}{2})$ there are infinitely many n, k with $C = \text{RS}[F_p, \langle \omega \rangle, k]$, $\langle \omega \rangle$ a smooth subgroup of order $n = 2^t$, prime $p \equiv 1 \pmod n$, $p < n^A$, and words $f = X^{rm}$, $g = X^{(r-1)m}$ such that at

$$\delta = (1 - \rho) - \frac{2}{K \log n} \quad (\eta = 2/s = 2/(K \log n) \text{ below capacity})$$

there are $\geq n^C$ scalars z with $\Delta(f + zg, C) \leq \delta$, yet $\Delta([f, g], C^2) > \delta$ — a direct proximity-gap and CA failure. [OUR-DERIVED, INFERRED] **Finite-field caveat (key number) — a *finite-field inferred ceiling*, not a proven endpoint:** the construction needs a prime $p \in [4^s, 8^s]$, so $s \leq \log_2 |F|/3$ and the gap is a *constant* $\eta \geq 6/\log_2 |F|$ (≈ 0.047 at 128-bit, ≈ 0.023 at 256-bit), **not** $o(1)$ at deployed sizes. The calculator encodes this as $\delta_{\text{unsafe}} = (1 - \rho) - 6/\log_2 |F|$. **This continuum endpoint is *inferred arithmetic* from the $8^s \leq |F|$ window — flagged VERIFY — and is retained only as an explicitly-labeled *pre-R13 continuum* asymptotic reference (the R13 power-of-two mechanism-asymptotic reference at 256-bit is $(1 - \rho) - 2/64$, §2.4), not a proven δ_{unsafe} ; the proven points underneath it are the positive-to-Johnson floor (radius-level, §2.1) and the unconditional BCHKS M31 witness at $\delta = \frac{1}{2}$ (over the non-smooth domain F_q^* — smooth-subgroup landing VERIFY; below). The per-field threshold version is the assembled Lemma below (R13): $\delta_C^* \leq (1 - \rho) - 2/s_{\max}(b)$ with $s_{\max}(b)$ a power of two. (Kambiré states it for prime F_p with $\rho \in (0, \frac{1}{2})$ open; the *mechanism* is established field-agnostic at EMF grade — see §3.2 — and the $\rho = \frac{1}{2}$ boundary is covered by the unconditional M31 BCHKS witness.)**

Lemma (per-field instantiated negative ceiling — ASSEMBLED, R13). Let $F = F_p$ be a prime field with $b = \log_2 p$ bits, let the smooth evaluation domain have order $n = 2^t$ with $n \mid p - 1$, and let $s = n/m$ be the Kambiré quotient parameter — necessarily a **power of two** (R13 integrality). Suppose (i) **distinctness:** $p > \varphi(s)^{\varphi(s)}$ (the KK25 distinctness hypothesis; cited, not re-proved here — 2026-06-11 update: the published KKH 2026/782 Lemma 1 now proves the distinctness statement in print for $1 \leq r \leq s/2$, printed bound $p > s^{s/2}$, count $2^r \binom{s/2}{r}$ — covering the 64- and 128-bit non- $\rho = \frac{1}{2}$ rows by statement, the 31-

bit row only via the proof’s intermediate bound $p > (2r)^{\varphi(s)}$; see `literature/notes/kkh-2026-782-verification.md`); (ii) **count**: $N(s, \rho) > 2^{b-128}$, where $N(s, \rho)$ is the per-rate bad-scalar count — Kambiré’s $N0_{\text{sum}}$ for $\rho < \frac{1}{2}$, and the proven antipodal count $3^{s/2}$ for $\rho = \frac{1}{2}$ (this case carries the **N1 conditionality flag**, §3.3); (iii) the construction’s standing distinctness/degree hypotheses; (iv) **radius/size**: s large enough that the firing radius stays strictly above Johnson, $(1 - \rho) - 2/s > J = 1 - \sqrt{\rho}$, i.e. $s > 2/(\sqrt{\rho} - \rho)$ — numerically $s \geq 16$ suffices for all four deployed rates (the threshold is ≤ 10.7 , attained at $\rho = \frac{1}{16}$) — together with Kambiré’s standing parameter hypotheses ($s = n/m$ with $s \mid n$, r in the construction’s admissible range). Without (iv), count + distinctness alone would nominally admit small- s “witnesses” at radii at or below Johnson, contradicting the proven positive side (§2.1); the lemma is asserted only in the near-capacity band above J . (All tabled s_{max} values, $16/16/32$, satisfy (iv).) Then $\varepsilon_{\text{mca}}(C, \delta) > 2^{-128}$ at $\delta = (1 - \rho) - 2/s$, and hence (using monotonicity of the failure event in δ , as Kambiré/CS25 state their ceilings)

$$\delta_C^* \leq (1 - \rho) - \frac{2}{s_{\text{max}}(b)},$$

where $s_{\text{max}}(b)$ is the largest power of two satisfying (i)+(ii): $s_{\text{max}} = 16$ at $b = 31$ — anchored at **BabyBear** ($p = 15 \cdot 2^{27} + 1$, 2-adicity 27), where the lemma applies verbatim; **literal M31 caveat**: $F_{2^{31}-1}$ has $p - 1 = 2(2^{30} - 1)$ with $v_2(p - 1) = 1$, so it has **no smooth domain of order ≥ 4** and the lemma is vacuous over literal M31 — the M31 figure refers to the deployed circle/ F_{p^2} smooth subgroup of order $p + 1 = 2^{31}$ via the §3.2 field-agnostic mechanism — so the M31 row is properly read as **31-bit characteristic / 62-bit alphabet** (F_{p^2}), not a literal 31-bit-field row (distinctness calibrated at the characteristic $p \approx 2^{31}$, giving the same $s_{\text{max}} = 16$; count threshold 2^{b-128} trivial at the alphabet’s $b = 62$), that row inheriting N2’s **ESTABLISHED-MODULO-FORMALIZATION** grade rather than the prime-field **ASSEMBLED** grade — $s_{\text{max}} = 16$ at $b = 64$, $s_{\text{max}} = 32$ at $b = 128$. At $b = 256$ **no valid s exists** (distinctness allows $s = 64$ but the count $3^{32} = 2^{51} \ll 2^{128}$; for $\rho < \frac{1}{2}$ the $N0_{\text{sum}}$ count is smaller still — `n2_prize_threshold.json` shows all four rates infeasible), so the mechanism establishes **no threshold ceiling** there. *Status*: **ASSEMBLED from cited components** — Kambiré *arXiv:2604.09724 Thm. 1* (bad-scalar count) + the distinctness lemma (now **proven in print** for the 64/128-bit non-N1 rows: **KKH 2026/782 Lemma 1**; the 31-bit row proof-level only; $\rho = \frac{1}{2}$ out of printed range) + exact threshold arithmetic (`experiments/small_rs_atlas/n2_prize_threshold.py`) + the

*R13 power-of-two integrality correction. Soft spots (narrowed 2026-06-11): the 31-bit row rests on KKH-782's proof-internal bound $p > (2r)^{\varphi(s)}$ rather than a printed statement; the $\rho = \frac{1}{2}$ case is N1-conditional ($r = s/2 + 2$ exceeds the printed $r \leq s/2$ range; KKH-782's appendix all- δ quotient construction supports the method asymptotically but does not supply the per-field row); the deg-convention ($N0_{\text{fire}}$ strict vs $N0_{\text{sum}}$ Kambiré-literal) is noted and tracked per row. (The continuum $(1 - \rho) - 6/\log_2 |F|$ of §2.4 is **not** part of this lemma — it is retained only as an explicitly-labeled asymptotic reference; the bracket's lower end J and the M31 $\delta = \frac{1}{2}$ point remain proven.)*

[**PROVEN, conditional**] (**BCHKS Thm. 1.13**). Prime multiplicative-subgroup RS at $\gamma = \delta_{\min} - \Theta(1/\log n)$ forces proximity loss $\Theta(1/\log n)$ — conditional on additive-combinatorics Conjecture 1.12. [**PROVEN, unconditional — over the non-smooth domain F_q^* ; smooth-domain landing VERIFY**] for M31, $\rho \approx \frac{1}{2}$: RS[$F_{2^{31}-1}, F_q^*, (q-1)/2 + 2/\log_2 q$] has f, g that are $\frac{1}{2}$ -close for *all* z yet CA-distance $\geq \frac{1}{2} + \frac{1}{62} \approx 0.516$ (loss ≈ 0.016) — a genuine failure at radius $\delta = \frac{1}{2}$ for a deployed prime at rate $\approx \frac{1}{2}$, **but not, as stated, a smooth-domain prize witness**: F_q^* has order $q-1 = 2(2^{30}-1)$, not a power of two, so it is not smooth in ABF's sense, and whether the witness lands on a strictly power-of-two subgroup is **VERIFY** (per `bchks.md`). Wherever this ledger cites “the unconditional M31 witness”, it is this F_q^* (non-smooth-domain) statement.

$\Rightarrow \delta_{\text{known-negative}}$ (**prime**): per-field $(1 - \rho) - 2/s_{\max}(b)$ per the assembled **Lemma above (R13; KK25 cited not re-proved, $\rho = \frac{1}{2}$ N1-conditional); the continuum $(1 - \rho) - 6/\log_2 |F|$ remains a *finite-field inferred ceiling* (inferred from the window arithmetic, VERIFY, asymptotic reference only; *not* a proven endpoint, though the M31 $\delta = \frac{1}{2}$ point below it is proven — over the non-smooth F_q^* domain, smooth landing VERIFY); the **same inferred ceiling holds for genuine odd-characteristic extension fields F_{p^e}** (the smooth-domain counterexample *mechanism* is field-agnostic — its distinct-bad-scalar count is a characteristic-zero cyclotomic invariant, verified exactly; see §3.2 and `n2-verdict.md`). Extensions buy no advantage; the negative forecloses the same constant- $\eta = 2/s$ near-capacity band for both.**

2.3 The barrier: BCHKS Theorem 1.9

[**PROVEN**] (**BCHKS Thm. 1.9 = ABF Thm. 5.2; all RS**). For $\gamma = \text{LDR}_{F,D,q}(\delta_{\min}) + 2/n$ (just past the list-decoding radius *for list size* q), there exist f, g with $\geq q/(2n)$ close line points but $\Delta([f, g], C^2) \geq \delta_{\min} - 1/n$, so $a/q \geq 1/(2n)$ independent of q . **Consequence**: improving MCA / proximity gaps beyond the Johnson radius for *any* RS code **requires** first improving its list-decoding radius (at list size q) beyond Johnson — a separately hard, in-general-false-for-some-RS, long-open problem.

The barrier is *conditional, not absolute* (line-decoding-analysis.md §2.3): it does not forbid $\delta_C^* > J$; it proves that any *small-error* ($\varepsilon_{\text{mca}} \leq \text{poly}(n)/q$) certificate of $\delta_C^* > J$ for smooth RS is itself a list-decoding-beyond-Johnson result. The only crack is **large-error (M)CA** (error a/q not $\ll 1/n$ but still $\leq 2^{-128}$), which both CS25 and CGHLL flag as the realistic target. This is exactly why the prize poses *both* challenges. (The list-size challenge’s escape from this barrier is the short paper’s §4.)

2.4 The per-rate and per-field tables (R13–R18) and the 256-bit status

$\delta_{\text{known-positive}} = J$ is a field-type-insensitive endpoint and **proven**; the negative side’s *mechanism* is field-agnostic (the smooth-domain counterexample is a characteristic-zero cyclotomic invariant — §3.2, n2-verdict.md), so prime and genuine odd-char extension share one bracket — endpoint-wise; this says nothing about positive solution strategies (§3.5). **The negative endpoint δ_{neg} below is given by the assembled Lemma of §2.2 (R13)** — $\delta_{\text{neg}} = (1 - \rho) - 2/s_{\text{max}}(b)$ with $s_{\text{max}}(b)$ the largest power-of-two s passing the KK25 distinctness calibration $p > \varphi(s)^{\varphi(s)}$ ($s_{\text{max}} = 32$ at $b = 128$); soft spots per the lemma’s status tag: KK25 cited not re-proved, $\rho = \frac{1}{2}$ N1-conditional. At $b = 256$ the mechanism is **infeasible at $\varepsilon^* = 2^{-128}$** (count $3^{32} = 2^{51} \ll 2^{128}$), so there is **no threshold-established 256-bit ceiling below R_{cap}** (the generic CS25/Elias ceiling at $\approx R_{\text{cap}}$ stands) and the 256-bit mechanism-asymptotic reference is $(1 - \rho) - 2/64$ (R13 power-of-two convention, $s = 64$ the largest distinctness-allowed power of two — the primary reference convention used throughout this ledger, matching the N1 figures of §3.3; the pre-R13 continuum $(1 - \rho) - 6/256$ is retained only as a secondary continuum reference) — inferred, VERIFY — genuinely open, outside the assembled lemma’s reach; *not* a proven endpoint; the proven negative point is the unconditional M31 $\delta = \frac{1}{2}$ witness (non-smooth F_q^* domain, smooth landing VERIFY — §2.2). Values from problem-ledger.md §5 and calculator/out/delta_star_tables.md (target $\varepsilon^* = 2^{-128}$, $n = 2^{20}$; the generated δ_{unsafe} column is field-agnostic in mechanism — prime and extension rows show the identical inferred ceiling, no “open” upper end; note the generated column still shows the old continuum $(1 - \rho) - 6/\log_2 |F|$ values and awaits regeneration per R13).

Per-rate table (prime *and* genuine odd-char extension; the endpoints are field-type insensitive):

ρ	J (positive, proven, radius- level)	δ_{neg} @128-bit (assembled Lemma §2.2, R13: $s_{\text{max}} = 32$)	δ_{neg} @256-bit (mechanism- asymptotic only, NOT threshold- established — R13 power-of-two convention $(1-\rho)-2/64$, $s = 64$; pre-R13 continuum $(1-\rho)-6/256$ in brackets)	gap @128-bit (prime <i>and</i> extension)
1/2	0.29289	0.43750	(0.46875, reference only; continuum 0.47656)	[0.293, 0.438]
1/4	0.50000	0.68750	(0.71875, reference only; continuum 0.72656)	[0.500, 0.688]
1/8	0.64645	0.81250	(0.84375, reference only; continuum 0.85156)	[0.646, 0.813]
1/16	0.75000	0.87500	(0.90625, reference only; continuum 0.91406)	[0.750, 0.875]

Per-field headline table at $\rho = \frac{1}{2}$ — split by domain type and evidence grade (R13). The two blocks are deliberately separate: prime-field smooth-subgroup rows carry the assembled Lemma’s ASSEMBLED grade; circle/extension *interpretation* rows carry N2’s ESTABLISHED-MODULO-FORMALIZATION grade and must never be quoted at ASSEMBLED grade.

(a) *Prime-field smooth-subgroup rows (assembled Lemma §2.2 applies verbatim — ASSEMBLED grade):*

field	b (bits)	$s_{\max}(b)$	ceiling $(1 - \rho) - 2/s_{\max}$ at $\rho = \frac{1}{2}$	grade
BabyBear ($p = 15 \cdot 2^{27} + 1$, 2-adicity 27)	31	16	0.375	ASSEMBLED (KK25 cited not re-proved; $\rho = \frac{1}{2}$ N1- conditional)
Goldilocks	64	16	0.375	ASSEMBLED (same soft spots)
128-bit prime	128	32	0.4375	ASSEMBLED (same soft spots)
256-bit	256	none valid	no threshold- established ceiling below R_{cap} (count $3^{32} = 2^{51} \ll 2^{128}$); mechanism- asymptotic reference 0.46875 only	OPEN below R_{cap} ; generic CS25/Elias ceiling at $\approx R_{\text{cap}}$ stands

(b) *Extension/circle-domain interpretation rows (N2 mechanism transfer — NOT theorem grade, NOT the assembled prime-field grade):*

field	b	$s_{\max}$	ceiling at $\rho = \frac{1}{2}$	grade
M31 (circle/ F_{p^2} reading; 31-bit characteristic / 62-bit alphabet)	31 (char) / 62 (alphabet)	16	0.375 — <i>via</i> <i>circle/F_{p^2}</i> <i>smooth</i> <i>subgroup of</i> <i>order</i> $p + 1 = 2^{31}$; $N2$ <i>ESTABLISHED-</i> <i>MODULO-</i> <i>FORMALIZATION</i> <i>grade, not</i> <i>theorem</i> <i>grade; literal</i> <i>M31 has</i> $v_2(p - 1) = 1$, <i>no smooth</i> <i>subgroup of</i> <i>order ≥ 4</i>	EMF (N2, §3.4; distinctness calibrated at the character- istic)

Radius-level caveat for all tables. The positive endpoint J is the radius theorem of §2.1; at 31/64/128-bit fields the single-code bound *cannot* certify $\varepsilon_{\text{mca}} \leq 2^{-128}$ at any radius (§2.1), so for those fields the table brackets the threshold radius δ_C^* — it is **not** a certified $\varepsilon^* = 2^{-128}$ two-sided bracket. Only at ≥ 256 -bit fields does the positive side carry a 2^{-128} certificate ($\delta_{\text{safe}} = 0.29251 \approx J$). At $\rho = \frac{1}{2}$, $n = 2^{20}$: 256-bit prime gives $\delta_{\text{safe}} = 0.29251 (\approx J)$; the query term alone needs $t = 257$ queries for 128 bits ($-\log_2(1 - \delta_{\text{safe}}) = 0.4992$, $128/0.4992 \approx 256.4$), while the calculator’s *composite* count is $t = 272$ — larger because at δ_{safe} the certified ε_{mca} already nearly saturates the 2^{-128} budget, so the query term must cover only the residual slack $2^{-128} - \varepsilon_{\text{mca}} - |\Lambda|/|F| \approx 2^{-135.8}$ ($272 \times 0.4992 \approx 135.8$ bits); a 256-bit extension gives the same δ_{safe} and the same δ_{unsafe} . M31/BabyBear/Goldilocks are **infeasible** for the single-code certificate (§2.1). The headline framing (64 $\rightarrow$ 128 bits at $\rho = \frac{1}{2}$; short paper §1) is precisely whether δ_C^* can be pushed above J inside these brackets.

Per-field corrections and the 256-bit status (internal revision rounds R13–R18; problem-ledger.md §5–§6 and the VERIFY register). This is the assembled per-field negative-ceiling assessment — R13 is the power-of-two integrality correction used throughout; R14–R18 are the 2026-06-10 rescue-path assessments. The R13 correction moves every previously tabulated per-field ceiling (values at $\rho = \frac{1}{2}$): at $b = 31$ the old continuum value 0.30645 was too *pessimistic* — the honest ceiling is $(1 - \rho) - 2/16 = 0.375$ ($s_{\max} = 16$ under the KK25 $\varphi(s)^{\varphi(s)}$ calibration; the direct $b = 31$ anchor is **BabyBear** — literal M31 has no smooth domain of order ≥ 4 , its 0.375 trans-

ferring via the deployed circle/ F_{p^2} smooth subgroup at N2’s ESTABLISHED-MODULO-FORMALIZATION grade, per the §2.2 Lemma’s M31 caveat — the M31 row is 31-bit *characteristic* / 62-bit *alphabet*; this re-scoping strengthens, not weakens, the withdrawal — the direction of the correction, window wider than previously claimed, is unchanged), so the M31 open window is $[0.293, 0.375]$ — a **radius-level** window: like every sub-256-bit field, M31 cannot certify $\varepsilon_{\text{mca}} \leq 2^{-128}$ from the single-code bound (§2.1), so this brackets the threshold radius, not a certified $\varepsilon^* = 2^{-128}$ quantity — and the earlier “**almost closed**” reading is **withdrawn**; at Goldilocks ($b = 64$) the old 0.40625 was too *optimistic* (corrected to 0.375, $s_{\text{max}} = 16$); at $b = 128$ the old 0.45313 falls to 0.4375 ($s_{\text{max}} = 32$).

The 256-bit rescue paths (R14–R18 assessments, ledger VERIFY register, 2026-06-10). At $b = 256$ the mechanism establishes **no threshold ceiling**, and the two rescue paths are assessed as follows. **(a) Per-prime distinctness certification — BLOCKED with current tools.** Full distinctness is not needed (only $\text{count}_p \geq 2^{128} = p^{1/2}$), but even under the most optimistic characteristic-zero count the required $s \approx 162$ means distinguishing $\sim \binom{81}{40} \approx 2^{77}$ subset sums by enumeration (computationally infeasible), and the height/norm route fails for a fixed prime ($|\text{Norm}(\Delta)| \leq (2s)^{\varphi(s)} \approx 2^{676} > 2^{256}$, so individual divisibility cannot be excluded, and no union bound over the $\sim 2^{154}$ difference pairs is available). **(b) BGK-type subgroup-sumset growth — OPEN, and not close** (literature/notes/subgroup-sumset-growth.md, surveyed 2026-06-10): the needed statement is that ~ 83 -fold subset sums of the order- s subgroup $\mu_s \subset F_p^*$ with $s \approx 324 = p^{0.033}$ cover $\geq p^{1/2} = 2^{128}$ values; the best published bounds give $\sim 2^{15}$ — an ≈ 113 -bit gap — with covering known only at $\sim 2^{63}$ summands (Cipra–Cochrane), and the $s \asymp \log p$ regime named open even for qualitative cancellation. (Scoping caveat: this negative verdict is specific to the *tiny-subgroup* regime of the 256-bit ceiling; the positive keystone P' lives at much larger subgroups, $|H| = n = p^{0.12} p^{0.65}$ deployed, and the verdict must not be over-propagated to it.) The computed feasibility boundary is $b^* \approx 140$ bits — a Kambiré-type prize-level witness exists only for $|F| \leq 2^{140}$ (experiments/small_rs_atlas/results/n2_prize_threshold.json; n2-verdict.md). The honest 256-bit status for a negative ceiling **below** R_{cap} (i.e., a Kambiré-type near-capacity, constant- η ceiling) is therefore **OPEN — needs a new idea**; the generic CS25/Elias ceiling at $\approx R_{\text{cap}}$ (§2.2) stands at every field size, 256-bit included; the per-field values above carry the assembled Lemma’s soft spots verbatim (KK25 cited not re-proved; $\rho = \frac{1}{2}$ N1-conditional; deg-convention tracked).

3. The line-decoding frontier: core-bet verdict and the sub-lemma close-outs

The positive MCA route relies on **line-decodability** (ABF Def. 4.20 / GG25 Def. 3.1): C is (δ, a, b) -line-decodable if whenever $\geq a$ scalars γ put $f_1 + \gamma f_2$ within δ of *some* codeword $U(\gamma)$, then $\geq b$ of those codewords lie on a single line $\{u_1 + \gamma u_2\}$ of codewords. **[PROVEN] (GG25 Thm. 3.5 = ABF Thm. 4.21):** if C is $(\delta, a, n + 1)$ -line-decodable then $\varepsilon_{\text{mca}}(C, \delta) \leq a/|F|$. This is the entire positive route: a line-decoding theorem with threshold $a \leq 2^{-128}|F|$ at radius δ certifies $\delta_C^* \geq \delta$. The exact conjectural target is **CGHLL26 (eprint 2026/532, App. A.5, Conjecture 2):** RS over any prime-field domain is line-decodable up to the list-decoding capacity radius R_{cap} (CGHLL’s “Elias radius”, stated in the large- q -approximation convention — §1) with threshold $a = \ell(\theta)n + o(n)$, $\ell(\theta) \leq c_1 2^{c_2 H(\rho)/\eta}$ ($c_1 = c_2 = 1$ calibrated to KK25); “the same is true if the alphabet is from an extension field of F_p ” — originally stated without evidence, and now given its first concrete (negative-side) support: the KK25 worst case extends verbatim to genuine F_{p^ϵ} and sits *at/below* the threshold a below R_{cap} , exactly as on primes (§3.2, **n2-verdict.md**), so the extension clause is **consistent-and-tight**, not refuted.

3.1 Core-bet verdict

[Assessment, from line-decoding-analysis.md §7] The *strong* (up-to-MDS-capacity, small-error) core bet is **DEAD**: CS25 Cor. 1 gives $\varepsilon_{\text{ca}} = 1$ below $1 - \rho$, and Kambiré gives explicit smooth-prime bad lines at $(1 - \rho) - \Theta(1/\log n)$. The achievable ceiling is the **list-decoding capacity radius** R_{cap} , short of MDS capacity by $\sim 1/\log_2 q$. Moreover Thm 1.9 collapses the positive *small-error* MCA route into an equally hard list-decoding problem (§2.3), so the realistic positive form is **large-error** line-decodability up to $R_{\text{cap}} - \Theta(1/\log_2 |F|)$, with error $\varepsilon_{\text{mca}} = 2^{O(H(\rho)/\eta)} n/|F|$ exponential in $1/\eta$ (CGHLL §5.5: $\approx 0.83\beta - 0.91\beta$ bits/query, not the full β). The pessimistic scenario $\delta_C^* = J$ is **forced in char-2** (BCHKS Cor. 1.7: a bad line at $\delta = J$ with $a \geq n^{2-\epsilon}$, via subspace polynomials) but **unsupported in odd-characteristic prime fields** (no known construction puts a smooth-prime bad line in $(J, (1 - \rho) - O(1/\log n))$; the subspace-polynomial mechanism has no prime analogue — CGHLL’s central thesis). So the core bet, properly stated as *large-error line-decodability up to R_{cap}* , is **likely TRUE for smooth prime RS** but **currently unproven**, reducing to the smooth-domain list-size sub-lemma P’ below.

3.2 The extension-field case is settled at ESTABLISHED-MODULO-FORMALIZATION grade: the counterexample mechanism is field-agnostic

The smooth-domain negatives are stated over prime fields (Kambiré: Linnik primes $p \equiv 1 \pmod n$; BCHKS Thm 1.13: prime/Mersenne; KK25 Lemma 9: multiplicative subgroup of a *prime* field), and the char-2 results (BCHKS Cor.

1.7) are the opposite extreme. The intermediate case — RS over a smooth subgroup of an *odd-characteristic extension field* F_{p^e} (M31², BabyBear/Goldilocks extensions — the actually-deployed FRI fields) — was once flagged as the untouched opening. **It is now settled at EMF grade (n2-verdict.md, exact computation): the smooth-domain counterexample mechanism is FIELD-AGNOSTIC.** The firing condition ($e_2 = 0$), the bad scalar ($\lambda = e_1$), and the distinct-bad-scalar count are **characteristic-zero cyclotomic invariants** — e_1, e_2 are reductions mod the characteristic p of fixed elements of $\mathbb{Z}[\xi_s]$, independent of the extension degree e — so genuine odd-char extensions $GF(p^e)$ realize the *same* count as the prime field $GF(p)$, for primes outside the finite bad set excluded by the KK25 distinctness calibration (verified exactly, 10/10 unsaturated cases; a decoder-free certified bad line exists on genuine $GF(31^2)$, $GF(127^2)$ at $\delta = \frac{1}{2}, \frac{3}{4}$). The “no additive subspace” hope does not buy extension fields any extra safety, because the obstruction here is **multiplicative** (cyclotomic subset sums), not additive, and it lives in F_{p^e} identically to F_p .

Three consequences (all **ESTABLISHED-MODULO-FORMALIZATION**, fail-closed): (i) the negative extends to genuine odd-char extensions at the *same* constant- $\eta = 2/s$ band as primes — extensions are **not** a haven; (ii) the construction is **sub-threshold at 256-bit** for every target rate (the distinct-bad-scalar count needs $s \gtrsim 512$ to exceed 2^{128} under the rigorous binomial count — even the most optimistic characteristic-zero count needs $s \approx 162$ (§2.4) — but the Kambiré distinctness/Linnik window needs $s \lesssim 85$ at 256-bit — empty intersection; field-bit boundary $b^* \approx 140$, a witness existing only for $|F| \leq 2^{140}$ where the threshold degenerates), identical for prime and extension; (iii) it does **not** refute CGHLL Conj. 2’s extension clause — the firing radius $\delta = \text{capacity} - 2/s$ sits *below* R_{cap} (since $\eta = 2/s > 1/\log_2 p$, with margin far exceeding the ≤ 0.0017 large- q -approximation gap of §1), and the count $\approx 2^{H(2\rho)/(2\eta)}$ is at/below the conjecture’s exponential threshold $a = 2^{H(\rho)/\eta} n$, so the extension clause is *consistent-and-tight* exactly as the prime clause. **There is therefore no field-type opening: extensions are neither a haven nor a special vulnerability.** The genuine remaining frontier is the *positive* keystone P’ (§3.5). Note the asymmetry: the **negative endpoint mechanism is field-agnostic** (this section), but the positive **P’** route should first be proved for **prime-field multiplicative subgroups**, with the extension-field analogue stated separately (open) — P’ is *not* itself field-agnostic.

3.3 Sub-lemma N1 (extend Kambiré to $\rho = \frac{1}{2}$) — CONDITIONAL CLOSURE modulo a named distinctness extension (KK25 Lemma 9 beyond its stated range)

Sub-lemma N1 (n1-rho-half.md). *Over a smooth $\langle \omega \rangle \subset F_p$, $n = 2^t$, the line $f = X^{rm}, g = X^{(r-1)m}$ at rate exactly $\rho = \frac{1}{2}$ ($r = s/2 + 2$) has $\geq n^C$ proximates at $\delta = \frac{1}{2} - \eta$, $\eta = 2/s = \Theta(1/\log n)$, yet $\Delta([f, g], C^2) > \delta$. Verdict — **conditional closure modulo a***

named distinctness extension (KK25 Lemma 9 beyond its stated range). The two **unconditional** halves go through *verbatim* at $\rho = \frac{1}{2}$: (i) the degree obstruction $\deg g = (r-1)m > k$, and (ii) the structural no-CA bound $S^* \leq (r-1)m < rm$ (both pure root-counting, rate- and field-blind). The bad-scalar count is supplied by the **PROVEN** closed form $|H^{(+r)}| = \sum_u \binom{s/2}{u} 2^u = 2^{\Theta(s)}$ (antipodal-pair / $\{-1, 0, +1\}$ -vector bijection, verified vs brute force at $s = 8, 16, 32$). Kambiré’s own count $\binom{s/2}{r}$ is **identically zero** at $\rho = \frac{1}{2}$ (since $r = s/2 + 2 > s/2$) and **KK25 Lemma 9 is out of range** ($r \leq \varphi(m)/2 = s/4$ is violated); the true count is *larger*, not absent. **Payoff (in Kambiré’s $\deg \leq k$ convention — under strict $\deg < k$ the construction fires not at exact rate $\frac{1}{2}$ but at $\frac{1}{2} + 1/n$, an $O(1/n)$ slack; see the residual below):** $\delta_C^*(\frac{1}{2}) \leq \frac{1}{2} - \Theta(1/\log n)$ — concretely ≤ 0.4375 at 128-bit; at 256-bit **not threshold-established** (R13: count $3^{32} = 2^{51} \ll 2^{128}$; 0.46875 mechanism-asymptotic reference only, power-of-two convention $s = 64$ — §2.4) — at the single most-deployed rate. **Field-agnostic corollary.** The count is a characteristic-zero cyclotomic invariant, so genuine odd-char extensions F_{p^e} give the *identical* count (25 distinct close scalars at the tested instance over $GF(p^2)$); the $\rho = \frac{1}{2}$ endpoint extends to the deployed extension case (**n2-verdict.md** mechanism). **The single residual (named).** A routine, numerically-airtight-but-not-yet-formalized strengthening of KK25 Lemma 9’s $\{-1, 0, +1\}$ -cyclotomic-independence / resultant-distinctness bound past its stated $r \leq \varphi(m)/2$ range to $r = s/2 + 2$ — the *for-all-n* distinctness-mod- p guarantee. The resultant criterion ($\text{Res}(\Phi_s, Q) = 0 \Leftrightarrow$ collision, $|\text{Res}| \leq (2r)^{s/2}$) is verified with **zero** counterexamples across $s \in \{8, 16, 32\}$, but Lemma 9 *as a black box* does not cover $\rho = \frac{1}{2}$. The result also uses Kambiré’s own $\deg \leq k$ convention (strict $\deg < k$ gives zero firing at rate exactly $\frac{1}{2}$; the construction lives at rate $\frac{1}{2} + 1/n$, an $O(1/n)$ slack vanishing asymptotically). **Honest prize framing.** This is a **near-capacity** endpoint, *not* a closure of the open band: it forecloses only $(0.4375, 0.5)$ at $\rho = \frac{1}{2}$ /128-bit (at 256-bit nothing is foreclosed at threshold — R13); it does **not** touch the open (J, R_{cap}) band ($= (0.293, 0.492)$ at 128-bit / $(0.293, 0.496)$ at 256-bit) and is **consistent with** (does not refute) CGHLL Conj. 2 ($\delta_{\text{unsafe}} < R_{\text{cap}}$, count at/below the threshold). It supplies a conditional negative endpoint at the most-deployed rate; the genuine frontier remains the positive keystone P’. **Never “unconditional” / “proven” tout court — it is CONDITIONAL on the named Lemma-9 extension.**

3.4 Sub-lemma N2 (extend Kambiré to an odd-characteristic extension field) — SETTLED (ESTABLISHED-MODULO-FORMALIZATION: field-agnostic, sub-threshold, Conj-2-consistent)

Sub-lemma N2. *Originally posed as “construct an analogous bad line over a smooth subgroup $\langle \omega \rangle$ of an odd-characteristic extension field F_{p^e} .” Settled by exact computation (`n2-verdict.md`) at EMF grade; no longer an open attack. **Close-out.** The counterexample mechanism is **field-agnostic**: the distinct-bad-scalar count is a characteristic-zero cyclotomic invariant (e_1, e_2 are reductions mod p of fixed elements of $\mathbb{Z}[\xi_s]$, independent of the extension degree e), so $GF(p^e)$ realizes the *same* count as $GF(p)$ (verified exactly, 10/10 unsaturated cases; the relevant obstruction is the *multiplicative* relation among the ζ^i , which behaves identically over F_p and F_{p^e}). The bad line **does extend** structurally to genuine extensions (decoder-free certificate at $\delta = \frac{1}{2}, \frac{3}{4}$ over $GF(31^2), GF(127^2)$), but it is **sub-threshold at 256-bit** (count needs $s \gtrsim 512$ to exceed 2^{128} , distinctness needs $s \lesssim 85$ — empty window; boundary $b^* \approx 140$) and it does **not** refute CGHLL Conj. 2’s extension clause ($\delta = \text{capacity} - 2/s < R_{\text{cap}}$, count $\approx 2^{H(2\rho)/(2\eta)}$ at/below the threshold $a = 2^{H(\rho)/\eta_n}$). **Net: extensions are neither a haven nor a special vulnerability — N2 is settled at EMF grade, field-agnostic.***

3.5 The positive keystone P’ and the recommended priority

Sub-lemma P’ (smooth-domain small list just beyond Johnson — the positive keystone). *For $L = \langle \omega \rangle \subset F_p$, $|L| = n$, rate ρ , there is $\varepsilon_0 = \varepsilon_0(\rho) > 0$ with*

$$|\Lambda(\text{RS}[F_p, L, \rho n], J + \varepsilon_0)| \leq O_\rho(1/\varepsilon_0) \quad (\text{constant in } n).$$

Conclusion / payoff: via the GG25 Lemma-8 bootstrap (CGHLL A.5) this supplies line-decodability, hence $\varepsilon_{\text{mca}}(\text{RS}[F_p, L, \rho n], J + \varepsilon_0) \leq a(n)/|F| \leq 2^{-128}$ — the **first provable smooth-domain MCA point strictly beyond Johnson**. It is *weaker than CGHLL Conj. 1* (only ε_0 beyond J , not all the way to R_{cap}), dodging the hardest part of the barrier. **Field scope (not field-agnostic).** P’ as stated is a **prime-field multiplicative-subgroup** statement ($L = \langle \omega \rangle \subset F_p^\times$): the char-2 analogue is *false at Johnson* (BCHKS Cor. 1.7, subspace polynomials), and the genuine odd-char **extension-field** analogue is **OPEN** (extensions contain F_p -subspaces, so the “no additive subspaces” lever the prime route relies on is unavailable) — it must be stated separately. So, unlike the negative endpoint mechanism (which *is* field-agnostic, §3.2), P’ is *not* field-agnostic.

Recommended priority (line-decoding-analysis.md §7): (1) **P' first** — the positive keystone (a **prime-field multiplicative-subgroup** statement; *not* field-agnostic — its extension-field analogue is a separate, open problem), now the *single* genuine open attack: it simultaneously advances the list-decoding challenge and unlocks the MCA line-decoding theorem. (N1 carries a **conditional closure modulo the named KK25 Lemma-9 extension** (n1-rho-half.md) — the negative $\rho = \frac{1}{2}$ endpoint; N2 is **settled at EMF grade**, field-agnostic — neither is a remaining open attack.) A first target band for P': $\delta = J + \varepsilon_0$ for the smallest ε_0 beating Johnson, at $\rho = \frac{1}{2}$, M31 (the prime-field case); the genuine odd-char extension analogue is stated separately and remains open.

4. List-size companion records

The list-size bracket itself — the budget trichotomy, the Proposition with parts (A)/(B)/(C), and the Thm 1.9 separation — is the short paper's §3–§4. The ledger records the supporting tool-state items.

Johnson-floor constant provenance. The floor constant $1/(2\eta\rho)$ is ABF Cor. 3.3 exactly, verified in `listsize-analysis.md`; the ledger's VERIFY register still carries a stale `pre-listsize-analysis.md` row — “List size $n \cdot m/\eta$... VERIFY constant + exponents” — which is superseded by that verification.

Resolver state (R29/M2 follow-up). The resolver's $\delta^*[\text{Cap}]$ column (insensitive to $n \in \{2^{16}, 2^{20}, 2^{24}\}$ and $m \in \{2, 3\}$ to within the last digit; formerly computed under the alphabet cap `_QCAP_LOG2 = 200` that made its 256-bit entries 200-bit values — the cap is now 320, and the resolver plus `listsize-analysis.md` §6/§8 have been regenerated to the true $q = 2^{256}$ values) reports the *finite- n volume-bound crossing* $\approx R_{\text{cap}}$. That finite- n crossing is **rate-dependent in sign** relative to the large- q approximation $1 - H_q(\rho)$ at $n = 2^{20}$, 256-bit: $\approx 4.1 \cdot 10^{-7}$ *above* it at $\rho = \frac{1}{2}$, and $\approx 2 \cdot 10^{-5}$ *below* it at $\rho \leq \frac{1}{4}$ (verified by direct computation; the old “sits below” reading was wrong at $\rho = \frac{1}{2}$). This finite- n crossing is the **conjectural upper reach**, not a proven δ^* — the proven object is the bracket $[J - o(1), R_{\text{cap}}]$ (short paper §3). The regenerated 256-bit large- q -approximation values are 0.49609/0.74683/0.87288/0.93618 at $\rho = \frac{1}{2}/\frac{1}{4}/\frac{1}{8}/\frac{1}{16}$; the exact crossings R_{cap} are 0.49609/0.74681/0.87285/0.93616, differing at the fifth decimal for $\rho \leq \frac{1}{4}$. The 192-bit entries were unaffected by the old cap ($192 < 200$).

What remains open in sub-problem 2 [CONJECTURED]. A single open input governs the upper reach to $\approx R_{\text{cap}}$ for $|F| > 2^{128}$: can the interleaved list be proven $\leq$ a small constant (or $\text{poly}(n)$, i.e. $\leq B^{1/m}$) in the *worst case* for *smooth-domain* RS at radii between $J \approx 0.293$ and $R_{\text{cap}} \approx 0.492$ (128-bit, $\rho = \frac{1}{2}$)? (It does **not** bear on the knife-edge $|F| = 2^{128}$ *value*, which is already pinned at the UD supremum, short paper Prop. (B): at $B = 1$ any list of size 2 breaks the budget, and size-2 lists exist just past UD by the MDS midpoint

argument — a small-but- > 1 list past J , which is what the conjecture would deliver, cannot move it.) This is the smooth-domain interleaved-list conjecture (`listsize-analysis.md` §4), equivalent to derandomizing the random-RS list bound (ABF Thm. 3.6 / [AGL24]: $|\Lambda(C, 1 - \rho - \eta)| \leq 2(1 - \rho - \eta)/\eta$ w.h.p. over a *random* domain) to the structured smooth domain — ABF open problem §7.9 / our sub-lemma P'. It is itself a *list-decoding* statement (not a proximity-gap statement), so it too dodges Thm 1.9; it is gated instead by the different (also hard) open problem of RS list-decoding on structured domains. Note this is precisely the bound `line-decoding-analysis.md` §3 flags as a theorem only for random domains; it is **not** currently available for smooth domains. The experiment supports the conjecture's strong (constant-list) form at tiny fields (§5); no proof exists for cryptographic fields, and the Wave-1 experiments cannot see the cryptographic-scale deep-hole regime.

5. Experimental corroboration

Workstream C (`experiments/small_rs_atlas/`) is **exact full-codeword enumeration** over small fields ($q \leq 257$, $n \leq 256$) — *corroborative, not cryptographic-scale*. Three findings bear on the analysis above.

(a) **Smooth $\approx$ random/full everywhere [EXPERIMENTAL]**. Across 144 matched subgroup-vs-random cells, the smooth-random differential in mean close-count per line is $+0.037$ (std 0.207) and in bad-line fraction is -0.0014 (std 0.021) — centered on zero, bounded by sampling noise. Smooth subgroups are statistically indistinguishable from random domains; structured inputs (rational, degree- $k+1$) produce *fewer* bad lines than uniform-random, not more. This supports the positive / derandomization route (§3.1).

(b) **Onsets [EXPERIMENTAL]**. Bad lines (high close-count, no CA) onset essentially *exactly at Johnson* (mean onset — Johnson: subgroup -0.001 , random -0.005), identically across domain types. But **interleaved lists stay singleton far past Johnson**: the max interleaved list size is exactly 1 from below Johnson, through Johnson, departing from 1 only at $\approx$ capacity $-O(1/n)$ (the `interleaved_onset.csv` data give onset — capacity ≈ -0.05 to -0.10 , i.e. $c \approx 1-2$ in capacity $-c/n$) and exploding *at* capacity. $m = 3$ matches $m = 2$ exactly. This directly corroborates the sub-problem-2 benign-list story (short paper §3): the dangerous radius for interleaved lists is capacity $-O(1/n)$, not Johnson.

(c) **The Kambiré mechanism is a genuine smooth-specific algebraic excess — near capacity only [EXPERIMENTAL]**. Reproducing the monomial line $f = X^a, g = X^b$ on a smooth prime subgroup: the cleanest cell is $p = 257$, $n = 64$, $k = 2$, $\delta = 0.90625$ (= capacity -0.0625 , capacity = 0.9688): **SMOOTH close-count = 32 = predicted sumset size, vs matched RANDOM = 7** (counterexample_kambire.csv, smooth row at $\delta = 0.90625$ has `close_count=32`, `predicted_sumset_size=32`; random row

has `close_count=7`). This is a real smooth-specific excess corroborating the sunset mechanism — but **near capacity only**; below capacity $-O(1/\log n)$ no smooth-specific bad line appears.

Caveats [stated honestly]. (1) Field sizes are tiny ($q \leq 257$); the $1/|F|$ floor is huge, so 2^{-128} -scale effects are invisible — these are *structural* statements, not quantitative bounds at cryptographic scale. (2) Sampling over words/lines/targets (each distance/list exact); a rare structured configuration could be missed. (3) **Crucially, the tested “extension” domains in *this Wave-1 battery* were $\text{GF}(2^m) = \text{char-2}$** , which has its *own* known negatives (subspace polynomials) — **NOT odd-characteristic F_{p^e}** . The odd-char extension case (N2) is **not** probed by this Wave-1 battery; it is settled separately by the exact cyclotomic computation of `n2-verdict.md` (§3.2, EMF grade), which establishes the counterexample mechanism is **field-agnostic** — genuine F_{p^e} behaves exactly like F_p , sub-threshold at 256-bit, Conj-2-consistent. (4) The per-line common-agreement maximum S^* is an exact proxy for CA on a line, not the protocol’s ε_{mca} ; bridging the two is the theory step.

6. Status register

Claim	Status	Basis
$\delta_{\text{known-positive}} = J$ for smooth RS (all fields)	PROVEN	Bordage–Chiesa 2025/2051 Thm. 9.2; BCHKS 2025/2055 Thm. 1.5
Bordage–Chiesa MCA error $(m+\frac{1}{2})^7 n^2 d / (3\rho^{3/2} F)$	PROVEN	2025/2051 Thm. 9.2
BCHKS Thm 1.5 constant C_ρ (linear- n bound)	OUR-DERIVED (VERIFY)	η -exponent 5 firm; C_ρ a placeholder

Claim	Status	Basis
$\delta_{\text{known-negative (prime)}}$ $= (1 - \rho) - 6/\log_2 F $ — the headline continuum endpoint	INFERRED (finite-field inferred ceiling; VERIFY) — inferred arithmetic from the $8^s \leq F $ window, retained as asymptotic reference only; <i>not</i> a proven endpoint. The per-field threshold ceiling $(1 - \rho) - 2/s_{\max}(b)$ is the assembled Lemma (§2.2, R13) — assembled from cited components; KK25 cited not re-proved, $\rho = \frac{1}{2}$ N1-conditional PROVEN over the non-smooth domain F_q^* ($q - 1 = 2(2^{30} - 1)$ not a power of two); smooth-subgroup landing VERIFY — <i>not</i> a proven smooth-domain prize witness	Kambiré arXiv:2604.09724 Thm. 1 (asymptotic, proven) + our window arithmetic (inferred); per-field lemma assembled R13 (experiments/small_rs_atlas/n2_prize_threshold. the M31 $\delta = \frac{1}{2}$ point (next row) is the proven negative witness
M31, $\rho \approx \frac{1}{2}$ unconditional CA failure at $\delta = \frac{1}{2}$ (loss ≈ 0.016)	$\rho = \frac{1}{2}$ N1-conditional PROVEN over the non-smooth domain F_q^* ($q - 1 = 2(2^{30} - 1)$ not a power of two); smooth-subgroup landing VERIFY — <i>not</i> a proven smooth-domain prize witness	BCHKS Thm. 1.13 (M31 instantiation); caveat per <code>bchks.md</code> and §2.2
Thm 1.9 barrier: MCA-beyond-Johnson $\Rightarrow$ RS list-decoding- beyond-Johnson	PROVEN	BCHKS Thm. 1.9 = ABF Thm. 5.2

Claim	Status	Basis
$\delta_{\text{known-negative}}$ for genuine odd-char extension fields $= (1 - \rho) - 6/\log_2 F $ (same band as prime)	ESTABLISHED-MODULO-FORMALIZATION (field-agnostic mechanism) — the <i>mechanism</i> /count is field-agnostic and exact; the continuum <i>endpoint value</i> is the same finite-field inferred ceiling as the prime row (VERIFY for the continuum, asymptotic reference only; the per-field values follow the same assembled Lemma of §2.2, R13)	n2-verdict.md (Theorem + VERIFICATION): bad-scalar count is a char-0 cyclotomic invariant (exact); extends to $GF(p^e)$, sub-threshold at 256-bit, does not refute CGHLL Conj. 2
N1: $\rho = \frac{1}{2}$ endpoint $\delta_C^*(\frac{1}{2}) \leq \frac{1}{2} - \Theta(1/\log n)$ ($\leq 0.4375@128$; $@256$ not threshold-established per R13, 0.46875 asymptotic reference) — in Kambiré's $\deg \leq k$ convention (strict $\deg < k$: zero firing at exact $\rho = \frac{1}{2}$; construction lives at $\rho = \frac{1}{2} + 1/n$) — prime and F_{p^e}	CONDITIONALLY ESTABLISHED — conditional closure modulo a named distinctness extension (KK25 Lemma 9 beyond its stated range)	n1-rho-half.md : no-CA + degree obstruction verbatim/unconditional at $\rho = \frac{1}{2}$; count $ H^{(+r)} = \sum_u \binom{s/2}{u} 2^u = 2^{\Theta(s)}$ proven (replaces Kambiré's vacuous $\binom{s/2}{r}$, KK25 Lem. 9 out of range). Residual: extension of Lem. 9's $\{-1, 0, +1\}$ -independence/resultant bound past $r \leq \varphi(m)/2$ (zero counterexamples, $s \leq 32$). Near-capacity band only; consistent with CGHLL Conj. 2

Claim	Status	Basis
Per-field negative ceilings $\delta_C^* \leq (1 - \rho) - 2/s_{\max}(b)$, $s_{\max} = 16/16/32$ at $b = 31/64/128$ — at $\rho = \frac{1}{2}$: 0.375@31-bit (anchored at BabyBear; literal M31 has no smooth domain of order ≥ 4 — value transfers via the deployed circle/ F_{p^2} subgroup, i.e. 31-bit characteristic / 62-bit alphabet , at N2's ESTABLISHED- MODULO- FORMALIZATION grade, §2.2 Lemma caveat; M31 “almost closed” withdrawn ; window $[0.293, 0.375]$, radius-level — sub-256-bit fields carry no 2^{-128} positive certificate, §2.1/§2.4), 0.375@64-bit, 0.4375@128-bit 256-bit negative ceiling below R_{cap} (Kambiré-type near-capacity, constant- η ; any known mechanism) at $\varepsilon^* = 2^{-128}$; feasibility boundary $b^* \approx 140$ bits	ASSEMBLED (KK25 cited not re-proved; $\rho = \frac{1}{2}$ N1-conditional; deg-convention tracked; the literal-M31 instantiation N2-grade per the §2.2 caveat; radius-level brackets per §2.4)	assembled Lemma §2.2 (R13): Kambiré Thm. 1 count + KK25 distinctness + experiments/small_rs_atlas/n2_prize_threshold. problem-ledger.md §5–§6
256-bit negative ceiling below R_{cap} (Kambiré-type near-capacity, constant- η ; any known mechanism) at $\varepsilon^* = 2^{-128}$; feasibility boundary $b^* \approx 140$ bits	OPEN — no threshold-established ceiling below R_{cap} (the generic CS25/Elias ceiling at $\approx R_{\text{cap}}$, §2.2, stands at every field size); rescue paths assessed: per-prime certification blocked ($\sim 2^{77}$ enumeration; norm bound $2^{676} > 2^{256}$), BGK-type growth open (≈ 113 -bit gap)	ledger VERIFY register (2026-06-10); literature/notes/subgroup-sumset-growth.md; experiments/small_rs_atlas/results/n2_prize_th

Claim	Status	Basis
List-size trichotomy; crossover at $ F = 2^{128}$	OUR-DERIVED	arithmetic on B ; Johnson + Elias composition (short paper §3)
$\delta_C^{*(2)} = (1 - \rho)/2$ at $ F = 2^{128}$ (binding/UD) — as a supremum ; with closed balls the endpoint is not attained (discreteness)	PROVEN	short paper Prop. (B); MDS/Singleton
$\delta_C^{*(2)} \geq J - o(1)$ for $ F > 2^{128}$ (proven floor; precisely $J - \eta_{\min}$, $\eta_{\min} = 1/(2\rho B^{1/m})$)	PROVEN	Johnson Cor. 3.3 + interleaving; smooth domains
$\delta_C^{*(2)} \leq R_{\text{cap}}(q, \rho) =$ $H_q^{-1}(1 - \rho)$ for $ F > 2^{128}$ (proven ceiling — the inverse-entropy crossing; the large- q approximation $1 - H_q(\rho)$ differs from it beyond the fifth decimal at some rates, §1) $\Rightarrow$ proven bracket	PROVEN	Elias/CS25 lower-bound ceiling (list $\geq q^{\Omega(n)} \gg B$ above the crossing)
$\delta_C^{*(2)} \in [J - o(1), R_{\text{cap}}]$ for $ F > 2^{128}$	PROVEN BRACKET	compose the two rows above; short paper §3
Upper reach $\delta_C^{*(2)} \approx R_{\text{cap}}$ for $ F > 2^{128}$	CONJECTURAL	needs worst-case large-list (list $\leq B^{1/m}$) smooth-RS list-decoding beyond Johnson; ABF §7.9 / P'; not gated by Thm 1.9 but open
Sub-problem 2 decouples from Thm 1.9, with a standard list-bound bracket (vs sub-problem 1's barrier-gated state)	OUR-DERIVED (rigorous)	short paper §4 (it is a pure list statement)
Smooth-domain interleaved list $\leq P(\rho, \eta)^m$ up to capacity $-\eta$	CONJECTURED	needs structured-domain list-decoding (ABF §7.9)

Claim	Status	Basis
CGHLL Conj. 2: prime/extension RS line-decodable to R_{cap}	CONJECTURED (extension clause now <i>consistent-and-tight</i> , not refuted)	CGHLL 2026/532 App. A.5; <code>n2-verdict.md</code> : the KK25 counterexample extends to genuine F_{p^e} and sits at/below the threshold a below R_{cap} — first concrete (negative-side) support for the extension clause
Smooth $\approx$ random; onset at Johnson; interleaved singleton to capacity $-O(1/n)$; $m=3 \equiv m=2$	EXPERIMENTAL	small fields $q \leq 257$, corroborative
$p = 257$ smooth close-count 32 vs random 7 at $\delta = 0.906$	EXPERIMENTAL	<code>counterexample_kambire.csv</code> ; near-capacity only

7. What this ledger does and does not establish

- **The companion short paper establishes** (rigorous, from cited Johnson + Elias bounds) the **proven two-sided bracket** $\delta_C^{*(2)} \in [J - o(1), R_{\text{cap}}]$ **for the interleaved list-size grand challenge at** $|F| > 2^{128}$ (proven floor $J - \eta_{\min}$ with $\eta_{\min} = 1/(2\rho B^{1/m}) \leq 2^{-29}$ deployed, written $J - o(1)$; proven ceiling the inverse-entropy crossing R_{cap} ; e.g. $[0.293, 0.496]$ at $\rho = \frac{1}{2}$, 256-bit), the collapse to the UD supremum $(1 - \rho)/2$ at $|F| = 2^{128}$ (endpoint not attained — discreteness) and vacuity below — and, most importantly, that this bracket **uses no new list-decoding and therefore dodges the Thm 1.9 barrier**, demonstrating that the two grand challenges are not equivalent and that sub-problem 2 is decoupled from Thm. 1.9 with a standard list-bound bracket, while sub-problem 1 stays barrier-gated. The upper reach to $\approx R_{\text{cap}}$ is **conjectural** (it requires an unproven worst-case large-list smooth-RS list-decoding theorem beyond Johnson; ABF §7.9 / P'), not established anywhere in this work.
- **This ledger does record (ASSEMBLED, with named soft spots)** the assembled per-field MCA negative-ceiling **assessment** $\delta_C^* \leq (1 - \rho) - 2/s_{\max}(b)$ — at $\rho = \frac{1}{2}$: 0.375 at 31- and 64-bit, 0.4375 at 128-bit (§2.2 Lemma, R13; KK25 cited not re-proved, $\rho = \frac{1}{2}$ N1-conditional) — including the **withdrawal of the M31 “almost closed” reading** (honest window $[0.293, 0.375]$, not $[0.293, 0.306]$ — a radius-level window, §2.4:

sub-256-bit fields carry no 2^{-128} positive certificate; M31 row read as 31-bit characteristic / 62-bit alphabet, §2.2); and it records that at **256-bit no threshold-established negative ceiling below R_{cap} — i.e., no Kambiré-type near-capacity, constant- η ceiling — exists from any known mechanism** (the generic CS25/Elias ceiling at $\approx R_{\text{cap}}$ stands at every field size), with both rescue paths assessed (per-prime certification blocked at $\sim 2^{77}$ enumeration; BGK-type growth ≈ 113 bits short) and the feasibility boundary $b^* \approx 140$ bits (§2.4).

- **Neither artifact claims** any solution to the **MCA grand challenge**. For sub-problem 1, only the *bracket* $[J, R_{\text{cap}})$ is known (bounded above near capacity by Kambiré/CS25/BCHKS, the negative mechanism established field-agnostic at EMF grade — the near-capacity negative holds for genuine odd-char extensions at the same band as prime, `n2-verdict.md`); a positive MCA point strictly above J for smooth RS is unproven and is exactly the prize.
- **What remains for the full prize:** (1) **MCA, positive (the genuine frontier):** prove sub-lemma P' (small smooth-domain list just past J) $\Rightarrow$ via GG25 bootstrap, the first $\delta_C^* > J$; the realistic ceiling is large-error line-decodability to $R_{\text{cap}} - \Theta(1/\log_2 |F|)$ (CGHLL Conj. 2), unproven for structured domains. **P' should first be proved for prime-field multiplicative subgroups; its extension-field analogue is stated separately and is open — P' is *not* field-agnostic** (only the negative endpoint mechanism is). (2) **MCA, negative:** the $\rho = \frac{1}{2}$ endpoint now carries a **conditional closure modulo a named distinctness extension (KK25 Lemma 9 beyond its stated range)** — sub-lemma N1 ($\delta_C^*(\frac{1}{2}) \leq \frac{1}{2} - \Theta(1/\log n)$, prime and F_{p^e} , **in Kambiré's $\deg \leq k$ convention:** under strict $\deg < k$ the construction fires at rate $\frac{1}{2} + 1/n$, not exact $\frac{1}{2}$ — §3.3); it forecloses only the near-capacity band, not the open (J, R_{cap}) . (N2, the odd-char extension negative, is **settled at EMF grade:** field-agnostic, sub-threshold at 256-bit, Conj-2-consistent — `n2-verdict.md`; no longer an open attack.) The genuinely open negative item is the **256-bit ceiling below R_{cap}** (Kambiré-type near-capacity, constant- η): no known mechanism reaches the 2^{-128} threshold below R_{cap} at $b = 256$ (the generic CS25/Elias ceiling at $\approx R_{\text{cap}}$ stands) — per-prime certification blocked at $\sim 2^{77}$ enumeration, BGK-type growth ≈ 113 bits short (§2.4) — with feasibility boundary $b^* \approx 140$ bits. (3) **List-size, upper reach:** the smooth-domain interleaved-list conjecture (CONJECTURED; ABF §7.9) — it governs only the conjectural reach to $\approx R_{\text{cap}}$ for $|F| > 2^{128}$; the knife-edge $|F| = 2^{128}$ *value* is pinned at the UD supremum (short paper Prop. (B)) and is not open.

Structural work toward the positive side is ongoing.

Acknowledgments

This work was carried out with substantial assistance from AI research tools — primarily Anthropic’s Claude, with independent adversarial review passes by OpenAI’s GPT models. All load-bearing claims are backed by exact-arithmetic verification scripts in the accompanying repository; the status tags reflect audited verification states, not model assertions.

References

(Exact eprint/arXiv ids from `literature/notes/`.)

- [ABF] G. Arnon, D. Boneh, G. Fenzi, *Open Problems in List Decoding and Correlated Agreement*, IACR eprint **2026/680** (Apr 8 2026). The survey framing the Proximity Prize; Defs. 2.9, 2.12, 4.1, 4.3; Cor. 3.3; Lemma 6.6; Thms. 4.12, 4.17, 4.21, 5.2, 5.3; open problem §7.9.
- [BCGM / Bordage–Chiesa] S. Bordage, A. Chiesa, Z. Guan, I. Manzur, *All Polynomial Generators Preserve Distance with Mutual Correlated Agreement*, IACR eprint **2025/2051** (rev. May 19 2026). Thm. 9.2 (positive MCA to Johnson, all polynomial generators); Lemma 10.1 (interleaving).
- [BCHKS] E. Ben-Sasson, D. Carmon, U. Haböck, S. Kopparty, S. Saraf, *On Proximity Gaps for Reed–Solomon Codes*, IACR eprint **2025/2055** (Nov 6 2025). Thm. 1.5 (positive, zero-loss to Johnson); Thm. 1.9 (the barrier); Cor. 1.7 (char-2 failure at J); Thm. 1.13 + Conj. 1.12 (prime mult.-subgroup; M31 unconditional).
- [CS25] E. Crites, A. Stewart, *On Reed–Solomon Proximity Gaps Conjectures*, IACR eprint **2025/2046** (Dec 19 2025). Thm. 1 / Lemma 1 / Cor. 1 (Elias volume bound; $\varepsilon_{\text{ca}} = 1$ strip); Thm. 7.4.1, Claim 1 (list-decoding capacity, $1/\log_2 q$ gap); Thm. 2 (CA $\Rightarrow$ list-decoding).
- [GG25] R. Goyal, V. Guruswami, *Optimal Proximity Gaps for Subspace-Design Codes and (Random) Reed–Solomon Codes*, IACR eprint **2025/2054** (rev. Mar 2026). Def. 3.1 (curve-decodability); Thm. 3.5 (line-decode $\Rightarrow$ MCA); Cors. 4.9/4.10, Thm. 5.16 (capacity for folded/subspace-design/random RS); Lemma 8 (the list-size bootstrap).
- [Kambiré] A. Kambiré, *Proximity Gaps Conjecture Fails Near Capacity over Prime Fields*, **arXiv:2604.09724** (Apr 2026). Thm. 1 (explicit smooth-prime counterexample at $(1 - \rho) - \Theta(1/\log n)$); the $f = X^{rm}$, $g = X^{(r-1)m}$ line; flesh-out of the Krachun–Kazanin (KK25) sketch.
- [CGHLL26] D. Carmon, L. Goldberg, U. Haböck, L. Lerer, I. Lesokhin (+ S. Papini, S. Samocha), *S-two Whitepaper*, IACR eprint **2026/532** (2026). App. A.5: the Elias radius (large- q approximation of R_{cap}); Conjecture 1 (list-decodability to it); **Conjecture 2** (line-decodability to it, the core bet, incl. extension-field clause); Thm. 37 (= KK25); Lemma 9 (distinct-sums count).

- **[Short paper]** Ö. Demirel, *Two Grand Challenges Are Not One*, `technical-note-short.md` (this repository, 2026-06-10). The polished companion: budget trichotomy, proven list-size bracket, Thm 1.9 separation, field-agnosticity of the negative mechanism, open problems.
- **Project artifacts (this work):** `problem-ledger.md` (formal bracket ledger); `listsize-analysis.md` + `calculator/listsize_resolution.py` (sub-problem-2 bracket analysis); `line-decoding-analysis.md` (core-bet verdict, sub-lemmas); `n1-rho-half.md` (sub-lemma N1 close-out: the $\rho = \frac{1}{2}$ negative endpoint, CONDITIONAL-CLOSURE modulo the named KK25 Lemma-9 extension, with the exact closed-form count); `n2-verdict.md` (sub-lemma N2 close-out: the field-agnostic extension case); `calculator/out/delta_star_tables.md` (machine-generated δ_C^* brackets; δ_{unsafe} column pre-R13, regeneration pending); `experiments/small_rs_atlas/` (exact enumeration, `counterexample_kambire.csv`, `interleaved_onset.csv`, `verify_rho_half_sums.py`); `experiments/small_rs_atlas/results/n2_prize_threshold.json` (the $b^* \approx 140$ feasibility boundary and per-field threshold arithmetic); `literature/notes/subgroup-sumset-growth.md` (the 256-bit rescue-path survey: BGK-type subgroup-sumset growth, ≈ 113 -bit gap).